

LDA final project

Kaitlin Maciejewski, Morgan de Ferrante, Bingnan Li, Volha Tryputsen

Contents

Exploratory	1
Smoking vs. Age, Sex	1
(1): Smoking ~ age, sex	1
Is there a relationship between age and smoking status?	1
Does this relationship differ by sex?	2
(2) number of cigarettes ~ age, sex	2
Is there a relationship between the number of cigarettes smoked per day and age?	2
All	2
Smokers only	2
Does this relationship differ by sex?	2
All	3
Smokers only	3
Smoking vs. health outcomes	3
(1) The relationship between current smoking status and systolic blood pressure.	3
smoking ~ sysbp	3
smoking ~ sysbp, sex	3
(2) The relationship between current smoking status and diastolic blood pressure.	4
smoking ~ diabp	4
smoking ~ diabp, sex	4
(3) The relationship between current smoking status and serum total cholesterol.	5
smoking ~ totchol	5
smoking ~ totchol, sex [!!!]	5
Cor plot	6
Missingness	7
Models	8
(1) Is there a relationship between age and smoking status? Does this relationship differ by sex? .	8
(2) Is there a relationship between the number of cigarettes smoked per day and age? Does this relationship differ by sex?	9
(3) Totchol and cursmoke	12
(4) Sysbp and cursmoke	14
(5) Diabp and cursmoke	16

Exploratory

Smoking vs. Age, Sex

(1): Smoking ~ age, sex

Is there a relationship between age and smoking status?

ANS: Yes, the proportion of smokers decreases with the age.

```
smoke_v_age = data %>%
  select(cursmoke, age) %>%
  ggplot(aes(x = age, y = cursmoke)) +
  geom_jitter(height = 0.1, alpha = 0.1) +
  geom_smooth(lwd = 1.5) +
  theme_bw()
```

Does this relationship differ by sex?

ANS: There is a higher proportion of smoker among men compared to women as both age ,but there is no interaction between age and sex.

```
smoke_age_sex = data %>%
  select(cursmoke, age, sex) %>%
  ggplot(aes(x = age, y = cursmoke, group = sex, color = sex)) +
  geom_jitter(height = 0.1, alpha = 0.1) +
  geom_smooth(lwd = 1.5) +
  theme_bw()
```

(2) number of cigarettes ~ age, sex

Is there a relationship between the number of cigarettes smoked per day and age?

ANS: Yes, number of sigarets smoked per day stays constant for 30-50 years old and decreases with age after 50 years old.

All

```
n_c_age_all = data %>%
  select(cigpday, age) %>%
  ggplot(aes(x = age, y = cigpday)) +
  geom_jitter(height = 0.1, alpha = 0.1) +
  geom_smooth(lwd = 1.5) +
  theme_bw()
```

Smokers only

```
n_c_age_smoke = data %>%
  select(cigpday, age) %>%
  filter(cigpday > 0) %>%
  ggplot(aes(x = age, y = cigpday)) +
  geom_jitter(height = 0.1, alpha = 0.1) +
  geom_smooth(lwd = 1.5) +
  theme_bw()
```

Does this relationship differ by sex?

ANS: There is sex effect (men smoke higer number of sigarets per day than women across age), but there is no sex and age interaction.

All

```
n_c_age_s_all = data %>%
  select(cigpday, age, sex) %>%
  ggplot(aes(x = age, y = cigpday, group = sex, color = sex)) +
  geom_jitter(height = 0.1, alpha = 0.1) +
  geom_smooth(lwd = 1.5) +
  theme_bw()
```

Smokers only

```
n_c_age_s_smoke = data %>%
  select(cigpday, age, sex) %>%
  filter(cigpday > 0) %>%
  ggplot(aes(x = age, y = cigpday, group = sex, color = sex)) +
  geom_jitter(height = 0.1, alpha = 0.1) +
  geom_smooth(lwd = 1.5) +
  theme_bw()
```

Smoking vs. health outcomes

(1) The relationship between current smoking status and systolic blood pressure.

smoking ~ sysbp

ANS: Proportion of smokers decreases with increase of systolic blood pressure

```
smoke_sbp = data %>%
  select(cursmoke, sysbp) %>%
  ggplot(aes(x = sysbp, y = cursmoke)) +
  geom_jitter(height = 0.1, alpha = 0.1) +
  geom_smooth(lwd = 1.5) +
  theme_bw()
```

ANS: slightly higher sysbp for non-smokers

```
smoke_sysbp_status = data %>%
  select(cursmoke, sysbp) %>%
  mutate(cursmoke = factor(cursmoke)) %>%
  ggplot(aes(y = sysbp, x = cursmoke)) +
  geom_boxplot(outlier.colour = "white") +
  theme_bw()
```

smoking ~ sysbp, sex

ANS: Proportion of smokers decreases with increase of systolic blood pressure; the proportion is higher for men (sex effect).

```
smoke_sysbp_sex = data %>%
  select(cursmoke, sysbp, sex) %>%
  ggplot(aes(x = sysbp, y = cursmoke, group = sex, color = sex)) +
  geom_jitter(height = 0.1, alpha = 0.1) +
  geom_smooth(lwd = 1.5) +
  theme_bw()
```

ANS: no differences in sysbp between male and female smokers and non-smokers

```
smoke_sysbp_sex_status = data %>%
  select(cursmoke, sex, sysbp) %>%
  mutate(cursmoke = factor(cursmoke),
         smoke_sex = interaction(cursmoke, sex)) %>%
  ggplot(aes(y = sysbp, x = smoke_sex)) +
  geom_boxplot(outlier.colour = "white") +
  theme_bw()
```

(2) The relationship between current smoking status and diastolic blood pressure.

smoking ~ diabp

ANS: Proportion of smokers decreases with increase of diastolic blood pressure for BP=100 and then proportion increases again (latter could be due to not enough data)

```
smoke_dbp = data %>%
  select(cursmoke, diabp) %>%
  ggplot(aes(x = diabp, y = cursmoke)) +
  geom_jitter(height = 0.1, alpha = 0.1) +
  geom_smooth(lwd = 1.5) +
  theme_bw()
```

ANS: no difference

```
smoke_dbp_box = data %>%
  select(cursmoke, diabp) %>%
  mutate(cursmoke = factor(cursmoke)) %>%
  ggplot(aes(y = diabp, x = cursmoke)) +
  geom_boxplot(outlier.colour = "white") +
  theme_bw()
```

smoking ~ diabp, sex

ANS: Proportion of smokers decreases with increase of diastolic blood pressure; the proportions are higher for men (sex effect).

```
smoke_dbp_s = data %>%
  select(cursmoke, diabp, sex) %>%
  ggplot(aes(x = diabp, y = cursmoke, group = sex, color = sex)) +
  geom_jitter(height = 0.1, alpha = 0.1) +
  geom_smooth(lwd = 1.5) +
  theme_bw()
```

ANS: no difference

```
smoke_dbp_s_bp = data %>%
  select(cursmoke, sex, diabp) %>%
  mutate(cursmoke = factor(cursmoke),
         smoke_sex = interaction(cursmoke, sex)) %>%
  ggplot(aes(y = diabp, x = smoke_sex)) +
  geom_boxplot(outlier.colour = "white") +
  theme_bw()
```

(3) The relationship between current smoking status and serum total cholesterol.

smoking ~ totchol

ANS: Proportion of smokers slightly decreases with increase of total cholesterol values

```
smoke_tc = data %>%
  select(cursmoke, totchol) %>%
  ggplot(aes(x = totchol, y = cursmoke)) +
  geom_jitter(height = 0.1, alpha = 0.1) +
  geom_smooth(lwd = 1.5) +
  theme_bw()
```

ANS: no difference

```
smoke_tc_bp = data %>%
  select(cursmoke, totchol) %>%
  mutate(cursmoke = factor(cursmoke)) %>%
  ggplot(aes(y = totchol, x = cursmoke)) +
  geom_boxplot(outlier.colour = "white") +
  theme_bw()
```

smoking ~ totchol, sex [!!!]

ANS: Proportion of smokers has nonlinear relationship with total cholesterol for women; proportions increases with increase in total cholesterol for men (sex by totchol interaction effect).

```
smoke_tc_sex = data %>%
  select(cursmoke, totchol, sex) %>%
  ggplot(aes(x = totchol, y = cursmoke, group = sex, color = sex)) +
  geom_jitter(height = 0.1, alpha = 0.1) +
  geom_smooth(lwd = 1.5) +
  theme_bw()
```

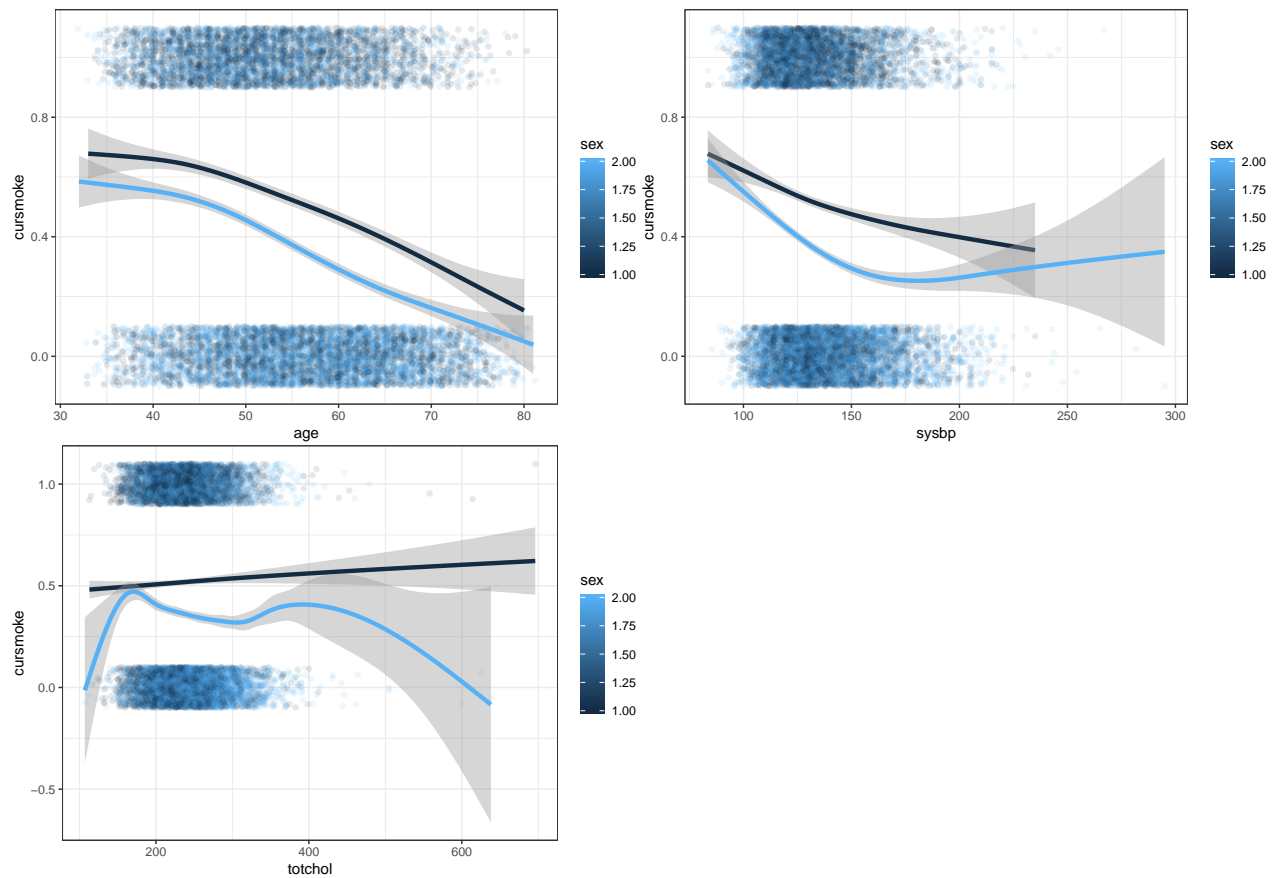
ANS: no difference

```
smoke_tc_sex_bp = data %>%
  select(cursmoke, sex, totchol) %>%
  mutate(cursmoke = factor(cursmoke),
         smoke_sex = interaction(cursmoke, sex)) %>%
  ggplot(aes(y = totchol, x = smoke_sex)) +
  geom_boxplot(outlier.colour = "white") +
  theme_bw()
```

```
smoke_age_sex
smoke_sysbp_sex
smoke_tc_sex
```

```
## Warning: Removed 409 rows containing non-finite values (stat_smooth).
```

```
## Warning: Removed 409 rows containing missing values (geom_point).
```




Cor plot

```
library(dplyr)

corr <- data[,c(-1,-21)] %>% cor(., use = "complete.obs")

library(corrplot)

corrplot(corr, type = "upper", order = "hclust",
          tl.col = "black", tl.srt = 75,
          tl.offset = 1, tl.cex = .8, method = "ellipse")
```



Missingness

```
prop <- round(colSums(is.na(data))/dim(data)[1], 3)

knitr::kable(sort(prop, decreasing = TRUE)[1:9], col.names = "Proportion of NAs")
```

Proportion of NAs	
hdlc	0.740
ldlc	0.740
glucose	0.124
bpmeds	0.051
totchol	0.035
educ	0.025
cigpday	0.007
bmi	0.004
hearttrte	0.001

```
prob.data <- data %>%
  group_by(period) %>%
  summarise(sysbp_prob = sum(sysbp, na.rm = TRUE)/n())
prob.data
```

```
## # A tibble: 3 x 2
##   period sysbp_prob
```

```
##      <int>      <dbl>
## 1         1       133.
## 2         2       137.
## 3         3       140.
```

```
table(data$period)
```

```
##
##      1      2      3
## 4434 3930 3263
```

Models

```
library(lme4)
library(dplyr)
```

(1) Is there a relationship between age and smoking status? Does this relationship differ by sex?

```
my.data <- read.csv("../final_data/frmgham2.csv")
library(gee)

model.q1 <- gee(CURSMOKE ~ AGE + as.factor(SEX) + as.factor(educ)
               + BMI + DIABETES + HEARTRTE + PREVCHD + PREVSTRK
               + PREVHYP + TIMEDTH,
               id = RANDID,
               data = my.data,
               family=binomial,
               corstr = "unstructured")
```

```
##      (Intercept)          AGE  as.factor(SEX)2  as.factor(educ)2
##      5.2810549757   -0.0570318456   -0.7246241605    0.0503157656
## as.factor(educ)3 as.factor(educ)4          BMI          DIABETES
##      -0.2263406725   -0.2104931713   -0.0940121564   -0.1114780579
##      HEARTRTE          PREVCHD          PREVSTRK          PREVHYP
##      0.0173911758   -0.0569530253   -0.2633153119   -0.2252617826
##      TIMEDTH
##      -0.0001021991
```

```
knitr::kable(summary(model.q1)$coefficients[,c(1,4,5)], digits = 3)
```

	Estimate	Robust S.E.	Robust z
(Intercept)	5.149	0.278	18.543
AGE	-0.051	0.002	-21.106
as.factor(SEX)2	-0.716	0.060	-11.892
as.factor(educ)2	0.083	0.071	1.168
as.factor(educ)3	-0.182	0.086	-2.116
as.factor(educ)4	-0.212	0.097	-2.185
BMI	-0.084	0.007	-12.152
DIABETES	-0.060	0.102	-0.595
HEARTRTE	0.009	0.001	6.031

	Estimate	Robust S.E.	Robust z
PREVCHD	-0.259	0.090	-2.872
PREVSTRK	-0.189	0.176	-1.073
PREVHYP	-0.049	0.041	-1.187
TIMEDTH	0.000	0.000	-6.341

```
model.q1[["working.correlation"]]
```

```
##           [,1]      [,2]      [,3]
## [1,] 1.0000000 0.7563191 0.5136670
## [2,] 0.7563191 1.0000000 0.5713635
## [3,] 0.5136670 0.5713635 1.0000000
```

(2) Is there a relationship between the number of cigarettes smoked per day and age? Does this relationship differ by sex?

```
model.q2 <- lmer(CIGPDAY ~ AGE + as.factor(SEX) +
  as.factor(educ) + (1|RANDID),
  data = my.data)
```

If we think cig per day as count data, it follows poisson distribution. Then we can fit GEE model as well:

```
model.q2_1 <- gee(CIGPDAY ~ AGE + as.factor(SEX) + as.factor(educ)
  + BMI + DIABETES + HEARTRTE + PREVCHD + PREVSTRK
  + PREVHYP + TIMEDTH,
  data = my.data,
  id = RANDID,
  family=poisson,
  corstr = "unstructured")
```

```
##      (Intercept)          AGE  as.factor(SEX)2  as.factor(educ)2
## 4.914149e+00 -3.395990e-02 -7.680017e-01 5.765596e-02
## as.factor(educ)3 as.factor(educ)4          BMI          DIABETES
## -1.227240e-01 -8.833837e-02 -4.493994e-02 -1.608903e-01
##      HEARTRTE      PREVCHD      PREVSTRK      PREVHYP
## 1.281640e-02 -9.883049e-02 -1.115583e-01 -1.079832e-01
##      TIMEDTH
## -5.152275e-05
```

```
knitr::kable(summary(model.q2_1)$coefficients[,c(1,4,5)], digits = 3)
```

	Estimate	Robust S.E.	Robust z
(Intercept)	4.523	0.170	26.626
AGE	-0.024	0.001	-15.826
as.factor(SEX)2	-0.795	0.039	-20.223
as.factor(educ)2	0.096	0.044	2.172
as.factor(educ)3	-0.072	0.057	-1.245
as.factor(educ)4	-0.067	0.063	-1.059
BMI	-0.039	0.005	-8.032
DIABETES	-0.110	0.072	-1.536
HEARTRTE	0.007	0.001	8.453
PREVCHD	-0.241	0.066	-3.635

	Estimate	Robust S.E.	Robust z
PREVSTRK	-0.130	0.110	-1.181
PREVHYP	-0.017	0.028	-0.587
TIMEDTH	0.000	0.000	-5.505

```
model.q2_1[["working.correlation"]]
```

```
##           [,1]      [,2]      [,3]
## [1,] 1.0000000 0.7425318 0.5006078
## [2,] 0.7425318 1.0000000 0.6615306
## [3,] 0.5006078 0.6615306 1.0000000
```

Using mixed effect model using cig per day instead of smoking status:

```
#saturated model
model.saturated <- lmer(CIGPDAY ~ as.factor(SEX) + AGE
                        + BPMEDS + as.factor(educ)
                        + TOTCHOL + BMI + GLUCOSE + DIABETES + HEARTRTE + PREVAP
                        + PREVCHD + PREVM I + PREVSTRK + STROKE + PREVHYP + (1|RANDID),
                        na.action = 'na.omit',
                        data = my.data)
summary(model.saturated)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula:
## CIGPDAY ~ as.factor(SEX) + AGE + BPMEDS + as.factor(educ) + TOTCHOL +
## BMI + GLUCOSE + DIABETES + HEARTRTE + PREVAP + PREVCHD +
## PREVM I + PREVSTRK + STROKE + PREVHYP + (1 | RANDID)
## Data: my.data
##
## REML criterion at convergence: 67692.5
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -4.4444 -0.3150 -0.1078  0.2067  6.2548
##
## Random effects:
## Groups   Name      Variance Std.Dev.
## RANDID   (Intercept) 95.78    9.787
## Residual                35.77    5.981
## Number of obs: 9310, groups:  RANDID, 4213
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)   23.237579   1.393589  16.675
## as.factor(SEX)2 -7.062484   0.339110 -20.827
## AGE          -0.178024   0.013735 -12.962
## BPMEDS         0.147068   0.340662   0.432
## as.factor(educ)2 0.602137   0.400578   1.503
## as.factor(educ)3 -0.984441   0.482907 -2.039
## as.factor(educ)4 -1.015103   0.548951 -1.849
## TOTCHOL        0.011407   0.002557   4.461
## BMI           -0.313257   0.035227 -8.893
## GLUCOSE       -0.010167   0.004203 -2.419
```

```
## DIABETES      -0.254321  0.570383 -0.446
## HEARTRTE      0.070622  0.008232  8.579
## PREVAP       -3.058426  0.985947 -3.102
## PREVCHD       0.948740  1.050262  0.903
## PREVMI       -2.594434  0.894896 -2.899
## PREVSTRK     -1.029740  0.917968 -1.122
## STROKE        0.943808  0.586941  1.608
## PREVHYP      -0.205972  0.240530 -0.856
```

#using variables that selected

```
model.mixed2 <- lmer(CIGPDAY~ AGE + as.factor(SEX) + SYSBP
                     + DIABP + TOTCHOL + as.factor(educ)
                     + (1|RANDID),
                     data = my.data,
                     na.action = "na.omit")
summary(model.mixed2)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula:
## CIGPDAY ~ AGE + as.factor(SEX) + SYSBP + DIABP + TOTCHOL + as.factor(educ) +
## (1 | RANDID)
## Data: my.data
##
## REML criterion at convergence: 78607.2
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -4.6343 -0.3025 -0.1015  0.1642  6.6280
##
## Random effects:
##   Groups   Name      Variance Std.Dev.
##  RANDID    (Intercept) 97.30    9.864
## Residual                36.52    6.043
## Number of obs: 10868, groups:  RANDID, 4306
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  21.750845  1.049612  20.723
## AGE          -0.199583  0.012148 -16.429
## as.factor(SEX)2 -6.448164  0.332812 -19.375
## SYSBP        -0.001364  0.006670  -0.205
## DIABP        -0.018833  0.011514  -1.636
## TOTCHOL       0.011959  0.002373   5.039
## as.factor(educ)2  1.055542  0.392720   2.688
## as.factor(educ)3 -0.625323  0.474072  -1.319
## as.factor(educ)4 -0.827192  0.539360  -1.534
##
## Correlation of Fixed Effects:
##              (Intr) AGE      a.(SEX SYSBP  DIABP  TOTCHO as.()2 as.()3
## AGE          -0.490
## as.fc(SEX)2 -0.144  0.019
## SYSBP        0.023 -0.460 -0.064
## DIABP        -0.452  0.314  0.076 -0.679
## TOTCHOL      -0.368 -0.095 -0.094  0.047 -0.154
## as.fctr(d)2 -0.225  0.110 -0.055  0.008  0.001  0.002
```

```
## as.fctr(d)3 -0.161  0.051 -0.092  0.028 -0.004 -0.007  0.351
## as.fctr(d)4 -0.175  0.044  0.074  0.032 -0.005 -0.002  0.301  0.244

library(dplyr)

#####
# Initial Models #
#####

# totchol_fit <- gee(totchol ~ cursmoke + age + factor(sex) + factor(educ) + bmi +
#                   diabetes + hearttrte + prevchd + prevhyp + prevstrk + death,
#                   id = randid,
#                   family = "gaussian",
#                   na.action = "na.omit")
#
# round(2 * pnorm(abs(coef(summary(totchol_fit))[,5]), lower.tail = FALSE), 3)
#
#
# sysbp_fit <- gee(sysbp ~ cursmoke + age + factor(sex) + factor(educ) + bmi +
#                  diabetes + hearttrte + prevchd + prevstrk + death,
#                  id = randid,
#                  family = "gaussian",
#                  na.action = "na.omit")
# round(2 * pnorm(abs(coef(summary(sysbp_fit))[,5]), lower.tail = FALSE), 3)
#
# diabp_fit <- gee(diabp ~ cursmoke + age + factor(sex) + factor(educ) + bmi +
#                  diabetes + hearttrte + prevchd + prevstrk + death,
#                  id = randid,
#                  family = "gaussian",
#                  na.action = "na.omit")
#
# round(2 * pnorm(abs(coef(summary(diabp_fit))[,5]), lower.tail = FALSE), 3)
```

(3) Totchol and cursmoke

```
#####
# Models After Removing Non Significant Terms #
#####

totchol_fit <- gee(totchol ~ cursmoke + age + factor(sex) + bmi +
                  diabetes + hearttrte + prevhyp ,
                  id = randid,
                  family = "gaussian",
                  corstr = "unstructured",
                  na.action = "na.omit")

## (Intercept)      cursmoke      age factor(sex)2      bmi
## 162.1032803    2.3672102    0.6268036    12.3710689    0.7814908
##      diabetes      hearttrte      prevhyp
##   -5.8194286    0.1891382    5.2800067

knitr::kable(summary(totchol_fit)$coefficients[,c(1,4,5)], digits = 3)
```

	Estimate	Robust S.E.	Robust z
(Intercept)	149.257	5.279	28.274
cursmoke	3.430	1.009	3.399
age	0.648	0.053	12.309
factor(sex)2	13.474	1.219	11.050
bmi	1.404	0.144	9.748
diabetes	-6.628	2.638	-2.513
heartrte	0.117	0.032	3.626
prevhyp	3.451	0.903	3.820

```
knitr::kable(round(2 * pnorm(abs(coef(summary(totchol_fit))[,5]), lower.tail = FALSE), 3))
```

	x
(Intercept)	0.000
cursmoke	0.001
age	0.000
factor(sex)2	0.000
bmi	0.000
diabetes	0.012
heartrte	0.000
prevhyp	0.000

```
QIC(totchol_fit)
```

```
##      QIC
## 84588.89
```

```
totchol_fit2 <- gee(totchol ~ cursmoke + age + factor(sex) + bmi +
  diabetes + heartrte + prevhyp ,
  id = randid,
  family = "gaussian",
  corstr = "exchangeable",
  na.action = "na.omit")
```

```
## (Intercept)      cursmoke      age factor(sex)2      bmi
## 162.1032803    2.3672102    0.6268036    12.3710689    0.7814908
##      diabetes      heartrte      prevhyp
## -5.8194286    0.1891382    5.2800067
```

```
knitr::kable(summary(totchol_fit2)$coefficients[,c(1,4,5)], digits = 3)
```

	Estimate	Robust S.E.	Robust z
(Intercept)	157.548	5.162	30.518
cursmoke	3.824	1.035	3.696
age	0.402	0.051	7.824
factor(sex)2	13.291	1.222	10.874
bmi	1.656	0.143	11.578
diabetes	-7.648	2.322	-3.293
heartrte	0.118	0.033	3.607
prevhyp	2.423	0.894	2.710

```
knitr::kable(round(2 * pnorm(abs(coef(summary(totchol_fit2))[,5]), lower.tail = FALSE), 3))
```

	x
(Intercept)	0.000
cursmoke	0.000
age	0.000
factor(sex)2	0.000
bmi	0.000
diabetes	0.001
hearttrte	0.000
prevhyp	0.007

```
QIC(totchol_fit2)
```

```
##      QIC
## 84588.27
```

(4) Sysbp and cursmoke

```
sysbp_fit <- gee(sysbp ~ cursmoke + age + factor(sex) + bmi +
  diabetes + hearttrte + prevchd + prevstrk + death,
  id = randid,
  family = "gaussian",
  corstr = "unstructured",
  na.action = "na.omit")
```

```
## (Intercept)      cursmoke      age factor(sex)2      bmi
## 40.1191453 -1.1246140 0.7236371 2.9078810 1.3070103
## diabetes hearttrte prevchd prevstrk death
## 4.8530466 0.2464253 2.4405511 8.6116353 7.0492511
```

```
knitr::kable(summary(sysbp_fit)$coefficients[,c(1,4,5)], digits = 3)
```

	Estimate	Robust S.E.	Robust z
(Intercept)	42.702	2.163	19.740
cursmoke	-0.702	0.433	-1.619
age	0.741	0.022	33.618
factor(sex)2	3.272	0.542	6.042
bmi	1.373	0.062	22.301
diabetes	3.664	1.127	3.250
hearttrte	0.175	0.016	11.231
prevchd	1.670	0.876	1.906
prevstrk	4.894	2.004	2.443
death	7.650	0.639	11.967

```
knitr::kable(round(2 * pnorm(abs(coef(summary(sysbp_fit))[,5]), lower.tail = FALSE), 3))
```

	x
(Intercept)	0.000
cursmoke	0.105

	x
age	0.000
factor(sex)2	0.000
bmi	0.000
diabetes	0.001
hearttrte	0.000
prevchd	0.057
prevstrk	0.015
death	0.000

```
QIC(sysbp_fit)
```

```
##      QIC
## 68837.26
```

```
sysbp_fit2 <- gee(sysbp ~ cursmoke + age + factor(sex) + bmi +
  diabetes + hearttrte + prevchd + prevstrk + death,
  id = randid,
  family = "gaussian",
  corstr = "exchangeable",
  na.action = "na.omit")
```

```
## (Intercept)    cursmoke      age factor(sex)2      bmi
## 40.1191453   -1.1246140   0.7236371   2.9078810   1.3070103
##      diabetes    hearttrte    prevchd    prevstrk    death
## 4.8530466    0.2464253    2.4405511    8.6116353    7.0492511
```

```
knitr::kable(summary(sysbp_fit2)$coefficients[,c(1,4,5)], digits = 3)
```

	Estimate	Robust S.E.	Robust z
(Intercept)	43.000	2.176	19.758
cursmoke	-0.772	0.431	-1.789
age	0.741	0.021	34.477
factor(sex)2	3.321	0.545	6.093
bmi	1.370	0.063	21.783
diabetes	3.173	1.114	2.848
hearttrte	0.172	0.016	10.967
prevchd	1.553	0.862	1.800
prevstrk	4.201	2.094	2.006
death	7.794	0.644	12.109

```
knitr::kable(round(2 * pnorm(abs(coef(summary(sysbp_fit2))[,5]), lower.tail = FALSE), 3))
```

	x
(Intercept)	0.000
cursmoke	0.074
age	0.000
factor(sex)2	0.000
bmi	0.000
diabetes	0.004
hearttrte	0.000
prevchd	0.072

	x
prevstrk	0.045
death	0.000

```
QIC(sysbp_fit2)
```

```
##      QIC
## 68837.55
```

(5) Diabp and cursmoke

```
diabp_fit <- gee(diabp ~ cursmoke + factor(sex) + factor(educ) + bmi +
  diabetes + hearttrte + prevstrk + death,
  id = randid,
  family = "gaussian",
  corstr = "unstructured",
  na.action = "na.omit")
```

```
##      (Intercept)      cursmoke factor(sex)2 factor(educ)2 factor(educ)3
##      47.3480563    -1.0968017    -0.8594048     0.7880155     0.6921165
## factor(educ)4          bmi      diabetes      hearttrte      prevstrk
##      0.5395935     0.9078065    -1.6162433     0.1550842     4.3067909
##           death
##      2.8662538
```

```
knitr::kable(summary(diabp_fit)$coefficients[,c(1,4,5)], digits = 3)
```

	Estimate	Robust S.E.	Robust z
(Intercept)	48.560	1.142	42.511
cursmoke	-0.607	0.243	-2.501
factor(sex)2	-0.591	0.295	-1.999
factor(educ)2	0.864	0.344	2.513
factor(educ)3	0.769	0.400	1.924
factor(educ)4	0.470	0.450	1.043
bmi	0.971	0.035	28.015
diabetes	-1.731	0.598	-2.895
hearttrte	0.112	0.009	12.846
prevstrk	1.842	1.207	1.526
death	3.219	0.336	9.582

```
knitr::kable(round(2 * pnorm(abs(coef(summary(diabp_fit))[,5]), lower.tail = FALSE), 3))
```

	x
(Intercept)	0.000
cursmoke	0.012
factor(sex)2	0.046
factor(educ)2	0.012
factor(educ)3	0.054
factor(educ)4	0.297
bmi	0.000

	x
diabetes	0.004
heartрте	0.000
prevstrk	0.127
death	0.000

```
QIC(diabp_fit)
```

```
##      QIC
## 53507.67
```

```
diabp_fit2 <- gee(diabp ~ cursmoke + factor(sex) + factor(educ) + bmi +
  diabetes + heartрте + prevstrk + death,
  id = randid,
  family = "gaussian",
  corstr = "exchangeable",
  na.action = "na.omit")
```

```
##      (Intercept)      cursmoke factor(sex)2 factor(educ)2 factor(educ)3
##      47.3480563    -1.0968017    -0.8594048     0.7880155     0.6921165
## factor(educ)4      bmi      diabetes      heartрте      prevstrk
##      0.5395935     0.9078065    -1.6162433     0.1550842     4.3067909
##      death
##      2.8662538
```

```
knitr::kable(summary(diabp_fit2)$coefficients[,c(1,4,5)], digits = 3)
```

	Estimate	Robust S.E.	Robust z
(Intercept)	48.526	1.156	41.980
cursmoke	-0.525	0.242	-2.166
factor(sex)2	-0.518	0.298	-1.738
factor(educ)2	0.857	0.346	2.476
factor(educ)3	0.776	0.403	1.927
factor(educ)4	0.531	0.453	1.171
bmi	0.991	0.035	28.079
diabetes	-2.072	0.582	-3.564
heartрте	0.106	0.009	12.102
prevstrk	1.049	1.215	0.863
death	3.317	0.338	9.800

```
knitr::kable(round(2 * pnorm(abs(coef(summary(diabp_fit2))[,5]), lower.tail = FALSE), 3))
```

	x
(Intercept)	0.000
cursmoke	0.030
factor(sex)2	0.082
factor(educ)2	0.013
factor(educ)3	0.054
factor(educ)4	0.241
bmi	0.000
diabetes	0.000
heartрте	0.000

	x
prevstrk	0.388
death	0.000

```
QIC(diabp_fit2)
```

```
##    QIC
## 53508
```